

## ✓ Customer Churn Prediction

### 1. Importing the dependencies

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import LabelEncoder
from imblearn.over_sampling import SMOTE
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
import pickle
```

### 2. Data Loading and Understanding

```
# load the csv data to a pandas dataframe
df = pd.read_csv("/content/WA_Fn-UseC_-Telco-Customer-Churn.csv")
```

```
df.shape
```

```
(7043, 21)
```

```
df.head()
```

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	OnlineSecurity
0	7590-VHVEG	Female	0	Yes	No	1	No	No phone service	DSL	No	No
1	5575-GNVDE	Male	0	No	No	34	Yes	No	DSL	Yes	Yes
2	3668-QPYBK	Male	0	No	No	2	Yes	No	DSL	Yes	Yes
3	7795-CFOCW	Male	0	No	No	45	No	No phone service	DSL	Yes	Yes
4	9237-HQITU	Female	0	No	No	2	Yes	No	Fiber optic	No	No

```
pd.set_option("display.max_columns", None)
```

```
df.head(2)
```

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	OnlineSecurity
0	7590-VHVEG	Female	0	Yes	No	1	No	No phone service	DSL	No	No
1	5575-GNVDE	Male	0	No	No	34	Yes	No	DSL	Yes	Yes

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
 #   Column              Non-Null Count  Dtype
---  -
 0   customerID          7043 non-null  object
 1   gender              7043 non-null  object
 2   SeniorCitizen       7043 non-null  int64
 3   Partner             7043 non-null  object
 4   Dependents          7043 non-null  object
 5   tenure              7043 non-null  int64
 6   PhoneService        7043 non-null  object
 7   MultipleLines        7043 non-null  object
 8   InternetService     7043 non-null  object
 9   OnlineSecurity      7043 non-null  object
```

```

10 OnlineBackup      7043 non-null object
11 DeviceProtection  7043 non-null object
12 TechSupport       7043 non-null object
13 StreamingTV       7043 non-null object
14 StreamingMovies   7043 non-null object
15 Contract          7043 non-null object
16 PaperlessBilling  7043 non-null object
17 PaymentMethod     7043 non-null object
18 MonthlyCharges    7043 non-null float64
19 TotalCharges      7043 non-null object
20 Churn             7043 non-null object
dtypes: float64(1), int64(2), object(18)
memory usage: 1.1+ MB

```

```

# dropping customerID column as this is not required for modelling
df = df.drop(columns=["customerID"])

```

```
df.head(2)
```

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	D
0	Female	0	Yes	No	1	No	No phone service	DSL	No	Yes	
1	Male	0	No	No	34	Yes	No	DSL	Yes	No	

```
df.columns
```

```

Index(['gender', 'SeniorCitizen', 'Partner', 'Dependents', 'tenure',
      'PhoneService', 'MultipleLines', 'InternetService', 'OnlineSecurity',
      'OnlineBackup', 'DeviceProtection', 'TechSupport', 'StreamingTV',
      'StreamingMovies', 'Contract', 'PaperlessBilling', 'PaymentMethod',
      'MonthlyCharges', 'TotalCharges', 'Churn'],
      dtype='object')

```

```
print(df["gender"].unique())
```

```
['Female' 'Male']
```

```
print(df["SeniorCitizen"].unique())
```

```
[0 1]
```

```
# printing the unique values in all the columns
```

```
numerical_features_list = ["tenure", "MonthlyCharges", "TotalCharges"]
```

```

for col in df.columns:
    if col not in numerical_features_list:
        print(col, df[col].unique())
        print("-"*50)

```

```

gender ['Female' 'Male']
-----
SeniorCitizen [0 1]
-----
Partner ['Yes' 'No']
-----
Dependents ['No' 'Yes']
-----
PhoneService ['No' 'Yes']
-----
MultipleLines ['No phone service' 'No' 'Yes']
-----
InternetService ['DSL' 'Fiber optic' 'No']
-----
OnlineSecurity ['No' 'Yes' 'No internet service']
-----
OnlineBackup ['Yes' 'No' 'No internet service']
-----
DeviceProtection ['No' 'Yes' 'No internet service']
-----
TechSupport ['No' 'Yes' 'No internet service']
-----
StreamingTV ['No' 'Yes' 'No internet service']
-----
StreamingMovies ['No' 'Yes' 'No internet service']
-----
Contract ['Month-to-month' 'One year' 'Two year']
-----
PaperlessBilling ['Yes' 'No']
-----

```

```
PaymentMethod ['Electronic check' 'Mailed check' 'Bank transfer (automatic)'
 'Credit card (automatic)']
-----
Churn ['No' 'Yes']
-----
```

```
print(df.isnull().sum())
```

```
gender      0
SeniorCitizen  0
Partner      0
Dependents   0
tenure       0
PhoneService  0
MultipleLines  0
InternetService  0
OnlineSecurity  0
OnlineBackup  0
DeviceProtection  0
TechSupport  0
StreamingTV   0
StreamingMovies  0
Contract      0
PaperlessBilling  0
PaymentMethod  0
MonthlyCharges  0
TotalCharges  0
Churn         0
dtype: int64
```

```
#df["TotalCharges"] = df["TotalCharges"].astype(float)
```

```
df[df["TotalCharges"]==" "]
```

```
gender SeniorCitizen Partner Dependents tenure PhoneService MultipleLines InternetService OnlineSecurity OnlineBackup
488 Female 0 Yes Yes 0 No No phone service DSL Yes No
753 Male 0 No Yes 0 Yes No No No internet service No internet service
936 Female 0 Yes Yes 0 Yes No DSL Yes Yes
1082 Male 0 Yes Yes 0 Yes Yes No No internet service No internet service
1340 Female 0 Yes Yes 0 No No phone service DSL Yes Yes
3331 Male 0 Yes Yes 0 Yes No No No internet service No internet service
3826 Male 0 Yes Yes 0 Yes Yes No No internet service No internet service
4380 Female 0 Yes Yes 0 Yes No No No internet service No internet service
5218 Male 0 Yes Yes 0 Yes No No No internet service No internet service
6670 Female 0 Yes Yes 0 Yes Yes DSL No Yes
6754 Male 0 No Yes 0 Yes Yes DSL Yes Yes
```

```
len(df[df["TotalCharges"]==" "])
```

```
11
```

```
df["TotalCharges"] = df["TotalCharges"].replace({" ": "0.0"})
```

```
df["TotalCharges"] = df["TotalCharges"].astype(float)
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 20 columns):
#   Column              Non-Null Count  Dtype
---  -
0   gender              7043 non-null   object
```

```
1 SeniorCitizen    7043 non-null    int64
2 Partner          7043 non-null    object
3 Dependents       7043 non-null    object
4 tenure           7043 non-null    int64
5 PhoneService     7043 non-null    object
6 MultipleLines    7043 non-null    object
7 InternetService  7043 non-null    object
8 OnlineSecurity   7043 non-null    object
9 OnlineBackup     7043 non-null    object
10 DeviceProtection 7043 non-null    object
11 TechSupport     7043 non-null    object
12 StreamingTV     7043 non-null    object
13 StreamingMovies  7043 non-null    object
14 Contract        7043 non-null    object
15 PaperlessBilling 7043 non-null    object
16 PaymentMethod   7043 non-null    object
17 MonthlyCharges  7043 non-null    float64
18 TotalCharges    7043 non-null    float64
19 Churn           7043 non-null    object
dtypes: float64(2), int64(2), object(16)
memory usage: 1.1+ MB
```

```
# checking the class distribution of target column
print(df["Churn"].value_counts())
```

```
Churn
No      5174
Yes     1869
Name: count, dtype: int64
```

### Insights:

1. Customer ID removed as it is not required for modelling
2. No missing values in the dataset
3. Empty strings in the TotalCharges column were replaced with 0
4. Class imbalance identified in the target

### 3. Exploratory Data Analysis (EDA)

```
df.shape
```

```
(7043, 20)
```

```
df.columns
```

```
Index(['gender', 'SeniorCitizen', 'Partner', 'Dependents', 'tenure',
       'PhoneService', 'MultipleLines', 'InternetService', 'OnlineSecurity',
       'OnlineBackup', 'DeviceProtection', 'TechSupport', 'StreamingTV',
       'StreamingMovies', 'Contract', 'PaperlessBilling', 'PaymentMethod',
       'MonthlyCharges', 'TotalCharges', 'Churn'],
      dtype='object')
```

```
df.head(2)
```

```
gender SeniorCitizen Partner Dependents tenure PhoneService MultipleLines InternetService OnlineSecurity OnlineBackup OnlineBackup
0 Female            0     Yes         No         1         No      No phone service          DSL              No              Yes
1 Male             0     No         No        34         Yes           No          DSL             Yes              No
```

```
df.describe()
```

	SeniorCitizen	tenure	MonthlyCharges	TotalCharges
count	7043.000000	7043.000000	7043.000000	7043.000000
mean	0.162147	32.371149	64.761692	2279.734304
std	0.368612	24.559481	30.090047	2266.794470
min	0.000000	0.000000	18.250000	0.000000
25%	0.000000	9.000000	35.500000	398.550000
50%	0.000000	29.000000	70.350000	1394.550000
75%	0.000000	55.000000	89.850000	3786.600000
max	1.000000	72.000000	118.750000	8684.800000

## Numerical Features - Analysis

Understand the distribution of the numerical features

```
def plot_histogram(df, column_name):

    plt.figure(figsize=(5, 3))
    sns.histplot(df[column_name], kde=True)
    plt.title(f"Distribution of {column_name}")


    # calculate the mean and median values for the columns
    col_mean = df[column_name].mean()
    col_median = df[column_name].median()

    # add vertical lines for mean and median
    plt.axvline(col_mean, color="red", linestyle="--", label="Mean")
    plt.axvline(col_median, color="green", linestyle="-", label="Median")


    plt.legend()

    plt.show()
```

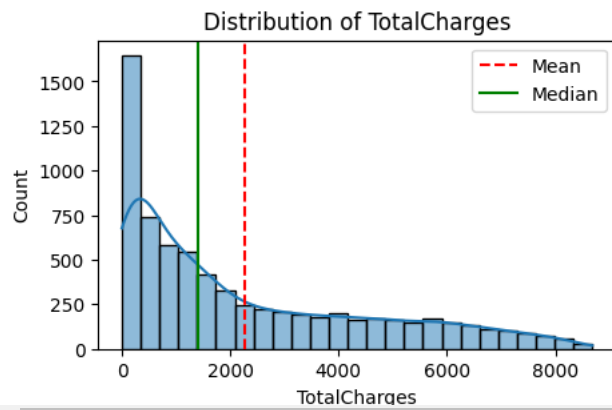
```
plot_histogram(df, "tenure")
```



```
plot_histogram(df, "MonthlyCharges")
```



```
plot_histogram(df, "TotalCharges")
```




### Box plot for numerical features


```
def plot_boxplot(df, column_name):
```

```
    plt.figure(figsize=(5, 3))
    sns.boxplot(y=df[column_name])
    plt.title(f"Box Plot of {column_name}")
    plt.ylabel(column_name)
    plt.show
```


```
plot_boxplot(df, "tenure")
```



```
plot_boxplot(df, "MonthlyCharges")
```




```
plot_boxplot(df, "TotalCharges")
```



### Correlation Heatmap for numerical columns

```
# correlation matrix - heatmap
plt.figure(figsize=(8, 4))
sns.heatmap(df[["tenure", "MonthlyCharges", "TotalCharges"]].corr(), annot=True, cmap="coolwarm", fmt=".2f")
plt.title("Correlation Heatmap")
plt.show()
```



### Categorical features - Analysis

```
df.columns
```

```
Index(['gender', 'SeniorCitizen', 'Partner', 'Dependents', 'tenure',  
      'PhoneService', 'MultipleLines', 'InternetService', 'OnlineSecurity',  
      'OnlineBackup', 'DeviceProtection', 'TechSupport', 'StreamingTV',  
      'StreamingMovies', 'Contract', 'PaperlessBilling', 'PaymentMethod',  
      'MonthlyCharges', 'TotalCharges', 'Churn'],  
      dtype='object')
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 7043 entries, 0 to 7042  
Data columns (total 20 columns):  
#   Column                Non-Null Count  Dtype  
---  ---  
0   gender                7043 non-null  object  
1   SeniorCitizen         7043 non-null  int64  
2   Partner               7043 non-null  object  
3   Dependents            7043 non-null  object  
4   tenure                7043 non-null  int64  
5   PhoneService          7043 non-null  object  
6   MultipleLines         7043 non-null  object  
7   InternetService       7043 non-null  object  
8   OnlineSecurity        7043 non-null  object  
9   OnlineBackup          7043 non-null  object  
10  DeviceProtection      7043 non-null  object  
11  TechSupport           7043 non-null  object  
12  StreamingTV           7043 non-null  object  
13  StreamingMovies       7043 non-null  object  
14  Contract              7043 non-null  object  
15  PaperlessBilling      7043 non-null  object  
16  PaymentMethod         7043 non-null  object  
17  MonthlyCharges        7043 non-null  float64
```

```

18 TotalCharges      7043 non-null   float64
19 Churn              7043 non-null   object
dtypes: float64(2), int64(2), object(16)
memory usage: 1.1+ MB

```

Countplot for categorical columns

```

object_cols = df.select_dtypes(include="object").columns.to_list()

object_cols = ["SeniorCitizen"] + object_cols


for col in object_cols:
    plt.figure(figsize=(5, 3))
    sns.countplot(x=df[col])
    plt.title(f"Count Plot of {col}")
    plt.show()

```

 [Show hidden output](#)

#### 4. Data Preprocessing


```
df.head(3)
```




	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	DeviceProtection
0	Female	0	Yes	No	1	No	No phone service	DSL	No	Yes	No
1	Male	0	No	No	34	Yes	No	DSL	Yes	No	No
2	Male	0	No	No	2	Yes	No	DSL	Yes	Yes	No

Label encoding of target column

```
df["Churn"] = df["Churn"].replace({"Yes": 1, "No": 0})
```


 `<ipython-input-57-b6eb27bc3ee0>:1: FutureWarning: Downcasting behavior in `replace` is deprecated and will be removed in a future version of pandas. To suppress this warning, please call `pandas.api.types.is_object_dtype(df["Churn"])` before calling `replace`. If you are using pandas < 1.0, you can avoid this warning by calling `df["Churn"] = df["Churn"].replace({"Yes": 1, "No": 0})`.`

```
df.head(3)
```



	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	DeviceProtection
0	Female	0	Yes	No	1	No	No phone service	DSL	No	Yes	No
1	Male	0	No	No	34	Yes	No	DSL	Yes	No	No
2	Male	0	No	No	2	Yes	No	DSL	Yes	Yes	No

```
print(df["Churn"].value_counts())
```

 Churn  
0 5174  
1 1869  
Name: churn, dtype: int64


Label encoding of categorical features

```

# identifying columns with object data type
object_columns = df.select_dtypes(include="object").columns

```

```
print(object_columns)
```

 Index(['gender', 'Partner', 'Dependents', 'PhoneService', 'MultipleLines', 'InternetService', 'OnlineSecurity', 'OnlineBackup', 'DeviceProtection', 'TechSupport', 'StreamingTV', 'StreamingMovies', 'Contract', 'PaperlessBilling', 'PaymentMethod'], dtype='object')



```
# initialize a dictionary to save the encoders
encoders = {}

# apply label encoding and store the encoders
for column in object_columns:
    label_encoder = LabelEncoder()
    df[column] = label_encoder.fit_transform(df[column])
    encoders[column] = label_encoder

# save the encoders to a pickle file
with open("encoders.pkl", "wb") as f:
    pickle.dump(encoders, f)
```

encoders

```
{'gender': LabelEncoder(),
 'Partner': LabelEncoder(),
 'Dependents': LabelEncoder(),
 'PhoneService': LabelEncoder(),
 'MultipleLines': LabelEncoder(),
 'InternetService': LabelEncoder(),
 'OnlineSecurity': LabelEncoder(),
 'OnlineBackup': LabelEncoder(),
 'DeviceProtection': LabelEncoder(),
 'TechSupport': LabelEncoder(),
 'StreamingTV': LabelEncoder(),
 'StreamingMovies': LabelEncoder(),
 'Contract': LabelEncoder(),
 'PaperlessBilling': LabelEncoder(),
 'PaymentMethod': LabelEncoder()}
```

df.head()

```
gender SeniorCitizen Partner Dependents tenure PhoneService MultipleLines InternetService OnlineSecurity OnlineBackup De
```

0	0	0	1	0	1	0	1	0	0	2
1	1	0	0	0	34	1	0	0	2	0
2	1	0	0	0	2	1	0	0	2	2
3	1	0	0	0	45	0	1	0	2	0
4	0	0	0	0	2	1	0	1	0	0

## Training and test data split

```
# splitting the features and target
X = df.drop(columns=["Churn"])
y = df["Churn"]
```

```
# split training and test data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
print(y_train.shape)
```

```
(5634,)
```

```
print(y_train.value_counts())
```

```
Churn
0    4138
1    1496
Name: count, dtype: int64
```

## Synthetic Minority Oversampling TEchnique (SMOTE)

```
smote = SMOTE(random_state=42)
```

```
X_train_smote, y_train_smote = smote.fit_resample(X_train, y_train)
```

```
print(y_train_smote.shape)
```

```
(8276,)
```

```
print(y_train_smote.value_counts())
```

```
↗ Churn
0    4138
1    4138
Name: churn, dtype: int64
```

## 5. Model Training

Training with default hyperparameters

```
# dictionary of models
models = {
    "Decision Tree": DecisionTreeClassifier(random_state=42),
    "Random Forest": RandomForestClassifier(random_state=42),
    "XGBoost": XGBClassifier(random_state=42)
}
```

```
# dictionary to store the cross validation results
cv_scores = {}

# perform 5-fold cross validation for each model
for model_name, model in models.items():
    print(f"Training {model_name} with default parameters")
    scores = cross_val_score(model, X_train_smote, y_train_smote, cv=5, scoring="accuracy")
    cv_scores[model_name] = scores
    print(f"{model_name} cross-validation accuracy: {np.mean(scores):.2f}")
    print("-"*70)
```

```
↗ Training Decision Tree with default parameters
Decision Tree cross-validation accuracy: 0.78
-----
Training Random Forest with default parameters
Random Forest cross-validation accuracy: 0.84
-----
Training XGBoost with default parameters
XGBoost cross-validation accuracy: 0.83
-----
```

cv\_scores

```
↗ {'Decision Tree': array([0.68297101, 0.71299094, 0.82175227, 0.83564955, 0.83564955]),
  'Random Forest': array([0.72524155, 0.77824773, 0.90513595, 0.89425982, 0.90090634]),
  'XGBoost': array([0.70048309, 0.75649547, 0.90271903, 0.89486405, 0.90030211])}
```

Random Forest gives the highest accuracy compared to other models with default parameters

```
rfc = RandomForestClassifier(random_state=42)
```

```
rfc.fit(X_train_smote, y_train_smote)
```

```
↗ RandomForestClassifier ⓘ ?
RandomForestClassifier(random_state=42)
```

```
print(y_test.value_counts())
```

```
↗ Churn
0    1036
1     373
Name: churn, dtype: int64
```

## 6. Model Evaluation

```
# evaluate on test data
y_test_pred = rfc.predict(X_test)

print("Accuracy Score:\n", accuracy_score(y_test, y_test_pred))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_test_pred))
print("Classification Report:\n", classification_report(y_test, y_test_pred))
```

```
↗ Accuracy Score:
0.7785663591199432
Confusion Matrix:
[[878 158]
```

```
[154 219]]
Classification Report:
              precision    recall  f1-score   support

     0       0.85        0.85        0.85        1036
     1       0.58        0.59        0.58         373

 accuracy          0.78          0.78          0.78        1409
 macro avg       0.72        0.72        0.72        1409
 weighted avg    0.78        0.78        0.78        1409
```

```
# save the trained model as a pickle file
model_data = {"model": rfc, "features_names": X.columns.tolist()}

with open("customer_churn_model.pkl", "wb") as f:
    pickle.dump(model_data, f)
```

## 7. Load the saved model and build a Predictive System

```
# load teh saved model and the feature names

with open("customer_churn_model.pkl", "rb") as f:
    model_data = pickle.load(f)

loaded_model = model_data["model"]
feature_names = model_data["features_names"]
```

```
print(loaded_model)
```

```
➡ RandomForestClassifier(random_state=42)
```

```
print(feature_names)
```

```
➡ ['gender', 'SeniorCitizen', 'Partner', 'Dependents', 'tenure', 'PhoneService', 'MultipleLines', 'InternetService', 'OnlineSecurity',
```

```
}

input_data = {
    'gender': 'Female',
    'SeniorCitizen': 0,
    'Partner': 'Yes',
    'Dependents': 'No',
    'tenure': 1,
    'PhoneService': 'No',
    'MultipleLines': 'No phone service',
    'InternetService': 'DSL',
    'OnlineSecurity': 'No',
    'OnlineBackup': 'Yes',
    'DeviceProtection': 'No',
    'TechSupport': 'No',
    'StreamingTV': 'No',
    'StreamingMovies': 'No',
    'Contract': 'Month-to-month',
    'PaperlessBilling': 'Yes',
    'PaymentMethod': 'Electronic check',
    'MonthlyCharges': 29.85,
    'TotalCharges': 29.85
}
```

```
input_data_df = pd.DataFrame([input_data])
```

```
with open("encoders.pkl", "rb") as f:
    encoders = pickle.load(f)
```

```
# encode categorical featires using teh saved encoders
for column, encoder in encoders.items():
    input_data_df[column] = encoder.transform(input_data_df[column])
```

```
# make a prediction
prediction = loaded_model.predict(input_data_df)
pred_prob = loaded_model.predict_proba(input_data_df)
```